

B+ TREE


By Abhishek RS

CONCEPT

- Used to implement Database indexes.
- B+ Trees are self-balancing, which means that as data is added or removed from the tree, it automatically adjusts itself to maintain a balanced structure.
- This ensures that the search time remains relatively constant, regardless of the size of the tree.
- In B+ tree the leaf nodes contains the actual data pointers.
- All leaf nodes remain at the same height.
- Leaf nodes are linked using Linked List.
- Therefore it supports both random and sequential access.

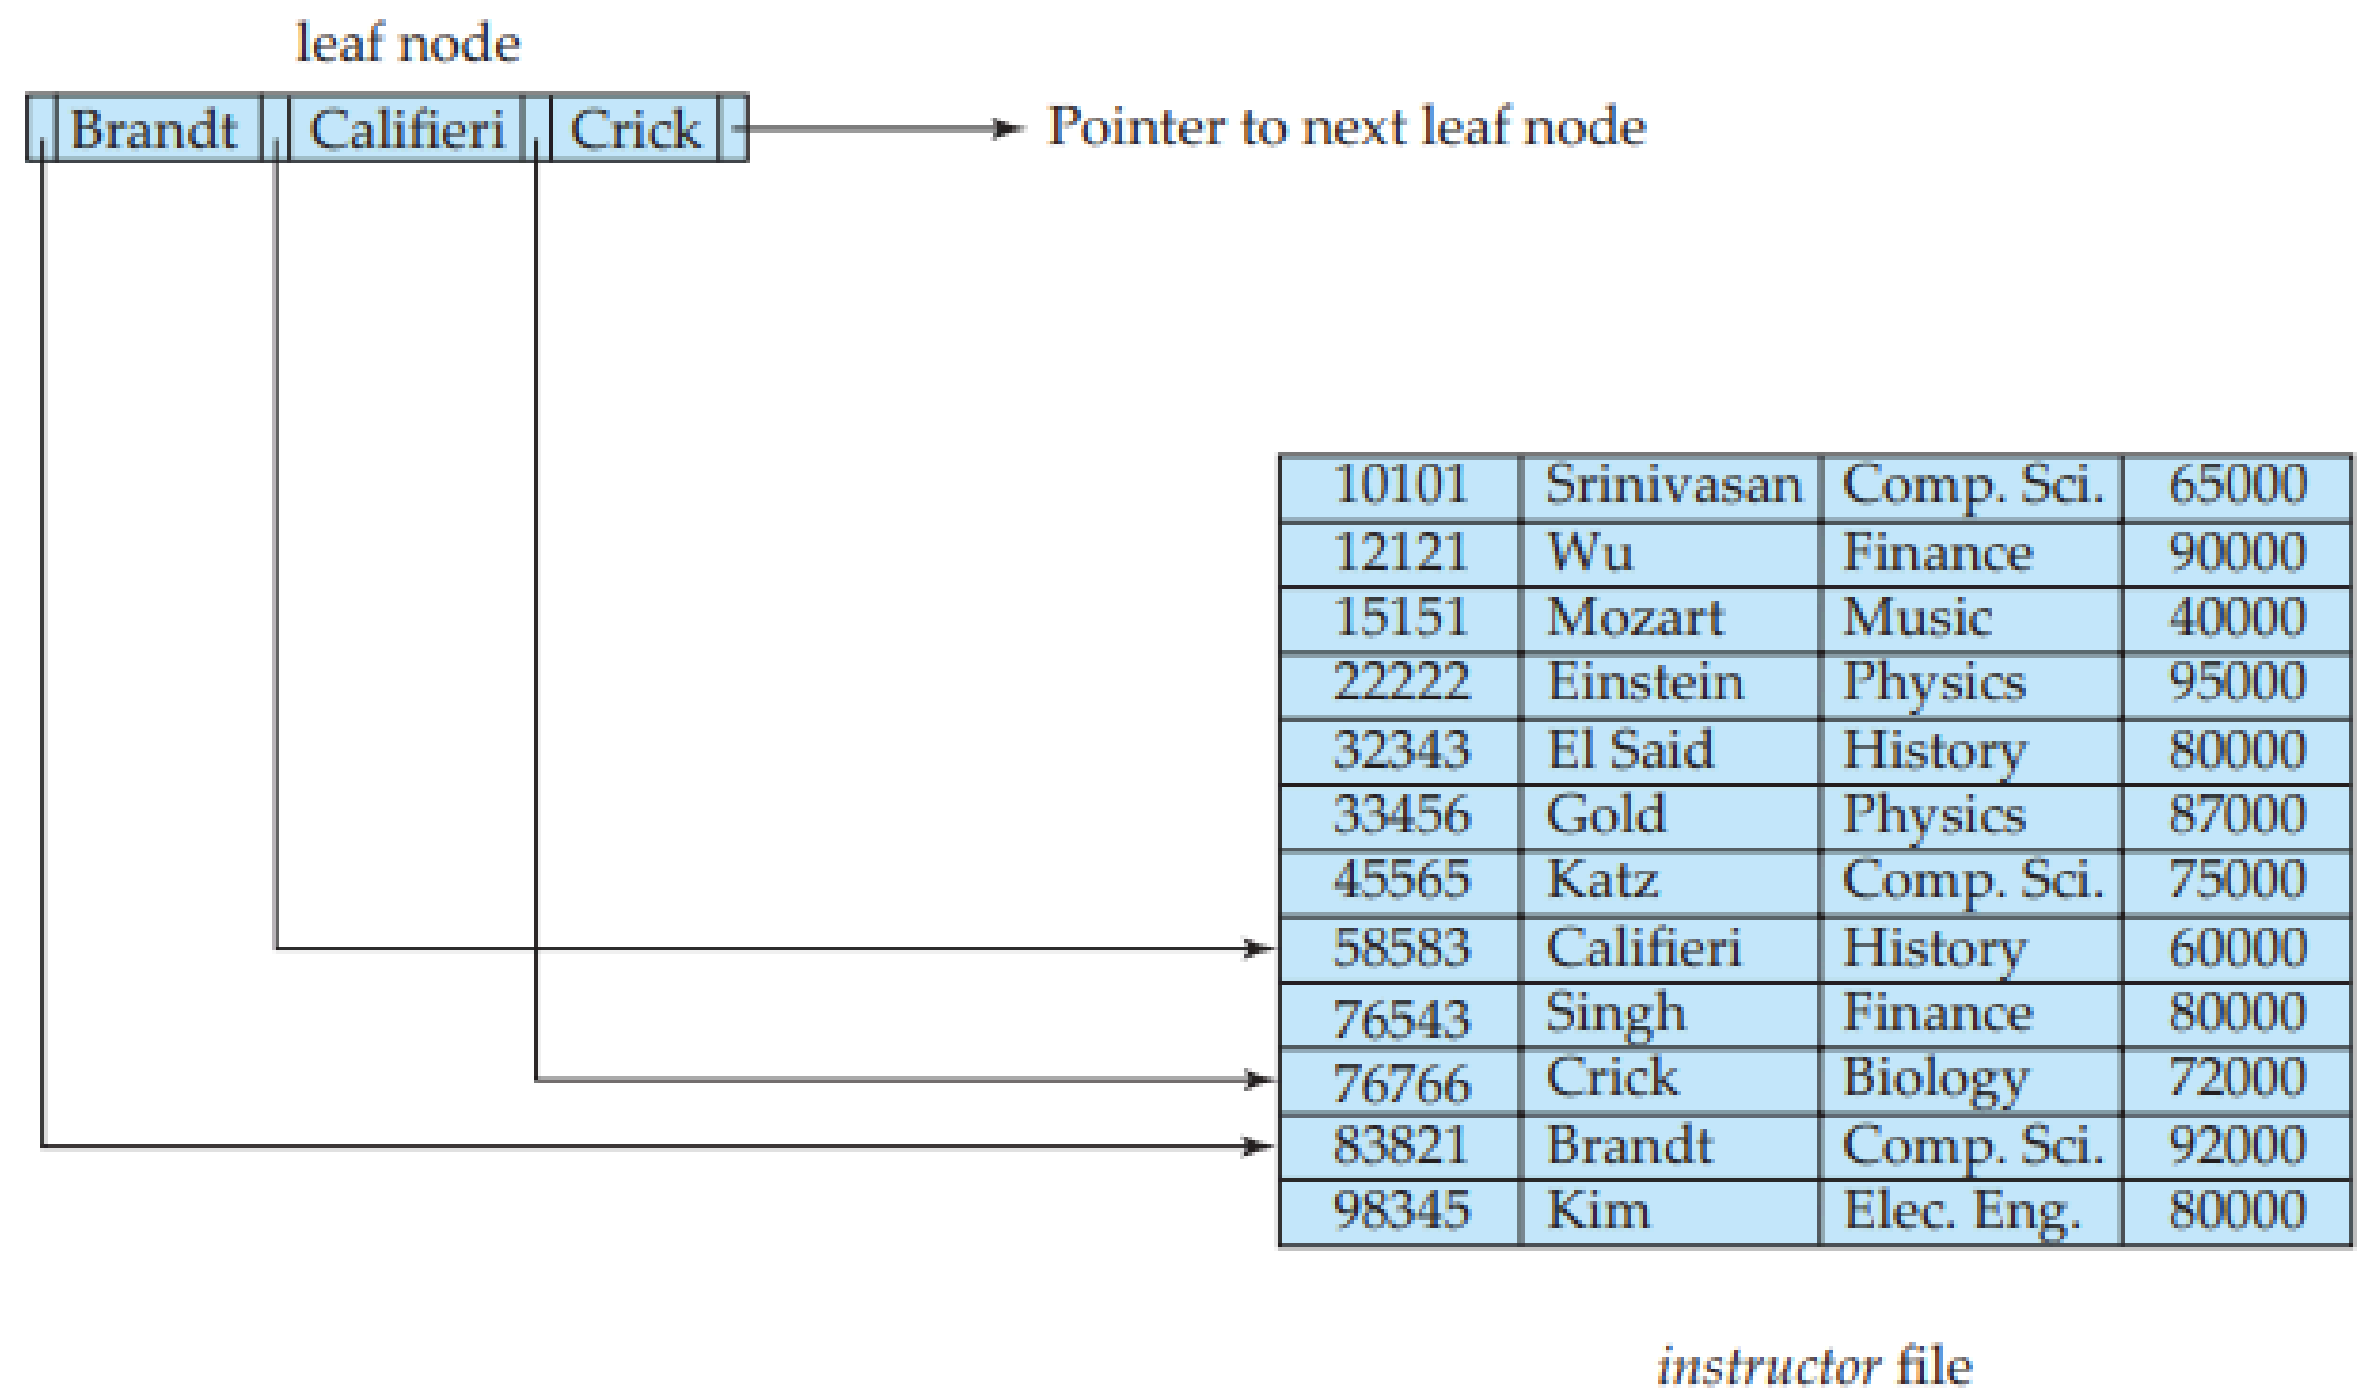
STRUCTURE OF B+ TREE

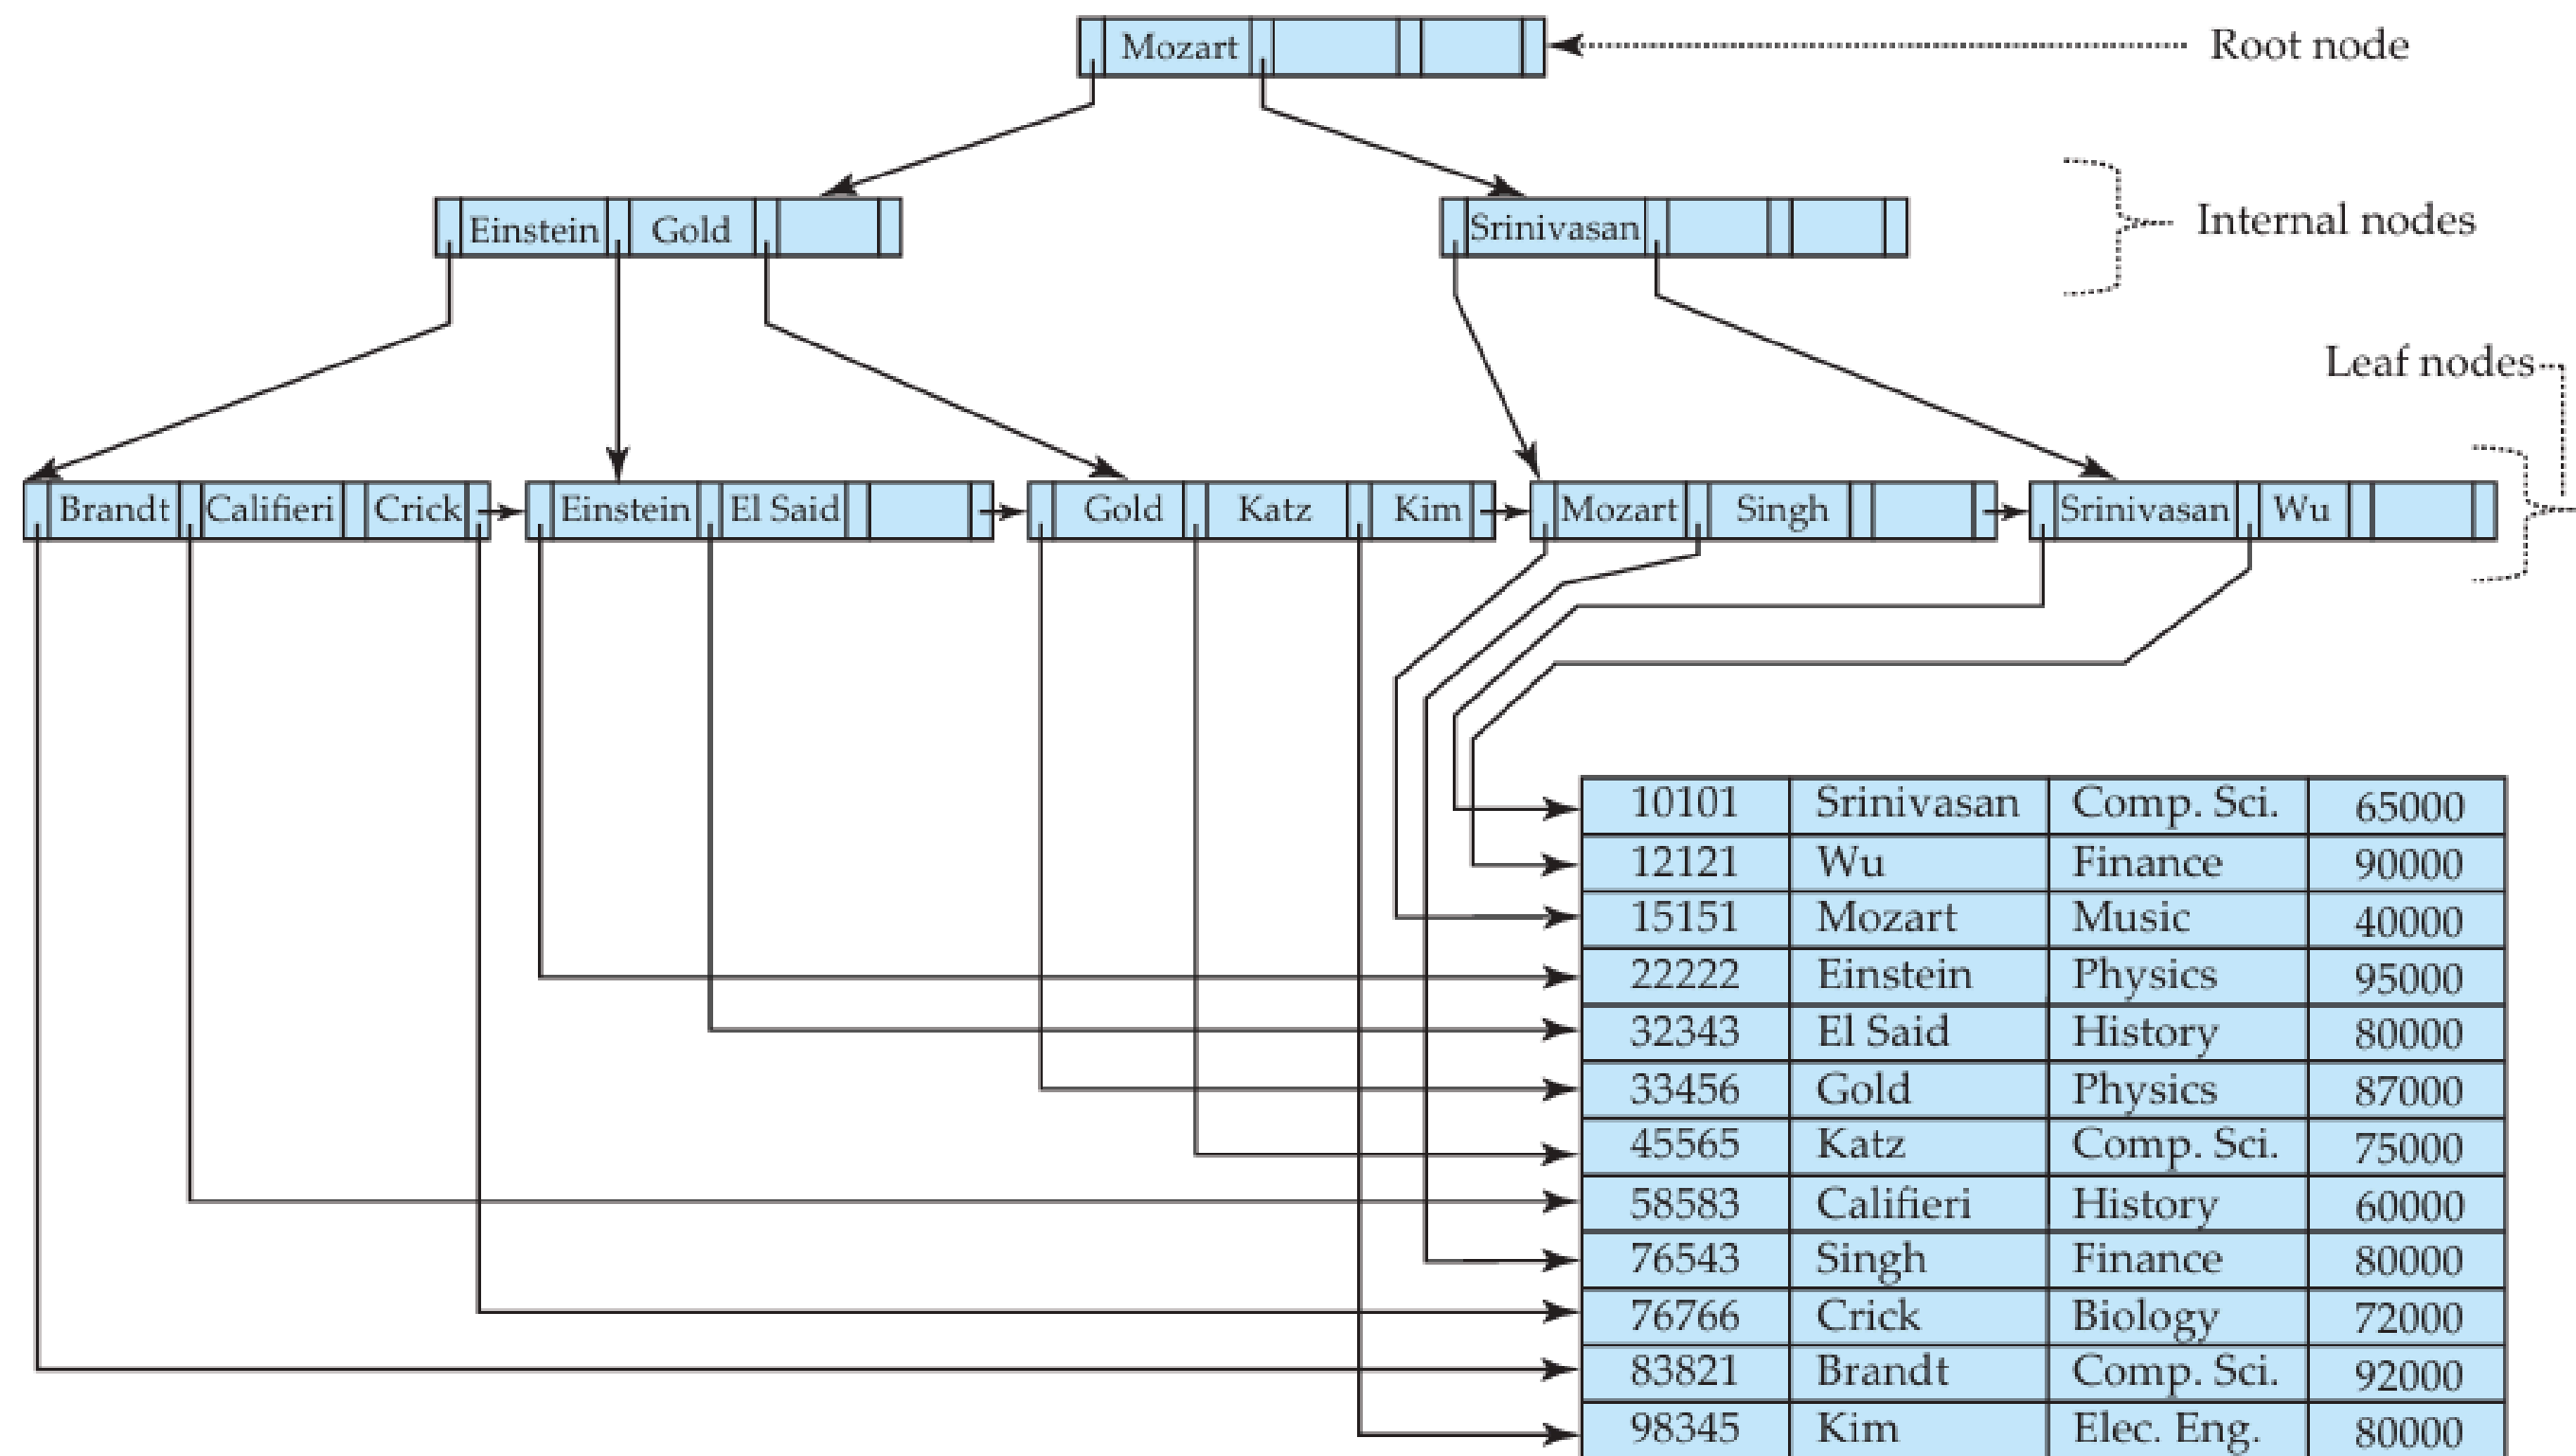
- A B+-tree index is a multilevel index, but it has a structure that differs from that of the multilevel index-sequential file



- This shows a typical node of a B+- tree. It contains up to $n - 1$ search-key values K_1, K_2, \dots, K_{n-1} , and n pointers P_1, P_2, \dots, P_n . The search-key values within a node are kept in sorted order; thus, if $i < j$, then $K_i < K_j$.

SIMPLE IMPLEMENTATION OF B+ TREE





ADVANTAGES

- Data stored in a B+ tree can be accessed both sequentially and directly.
- It takes an equal number of disk accesses to fetch records.
- A B+ tree with 'l' levels can store more entries in its internal nodes compared to a B-tree having the same 'l' levels.

DISADVANTAGES

- Because of internal nodes storing multiple keys and pointers, B+ Trees consume more memory compared to simpler indexing methods like hashing.
- Maintaining the balance of the tree requires restructuring operations (like node splitting/merging), which may slow down performance in write-heavy databases.
- Though searching in a B+ Tree is efficient ($O(\log n)$), accessing actual data requires traversing from the root to the leaf node, which may add overhead.



THANK YOU